

Distributed and Asynchronous Algorithm for Smooth High-dimensional Function Approximation using Orthotope B-splines

J. Meyer · C.C. de Visser

Abstract Aircraft are complex systems with, in some cases, high-dimensional non-linear interactions between control surfaces. When a failure occurs, adaptive flight control methods can be utilised to stabilise and make the aircraft controllable. Adaptive flight control methods, however, require accurate aerodynamic models - where first-order continuity is necessary for estimating the control derivatives and mitigating chattering that can reduce the longevity of components. Additionally, high-dimensional offline model identification with current approaches can take several hours for a few dimensions and this means model iterating and hyper-parameter tuning is often not feasible. Current approaches to smooth high-dimensional functional approximation are not scalable, require global communication between iteration steps, and are ill-conditioned in higher dimensions. This research develops the Distributed Asynchronous B-spline (DAB) algorithm that is more robust to ill-conditioning, due to low data coverage, by using first-order methods with acceleration and weighted constraint application. This algorithm is also suitable for continuous state-spaces. Smooth aerodynamic models can be determined in exactly $n \cdot \hat{r}$ iterations, where \hat{r} is the number of continuity equations in a single dimension and n is the number of dimensions. Moreover, memory reorganisation is proposed to avoid false sharing and conflict-free use of shared memory on the GPU to ensure that the algorithm runs efficiently in parallel.

Keywords Smooth High-dimensional Linear Function Approximation · Linearly-constrained Convex Quadratic Programming · Asynchronous Algorithms · Parallel Algorithms · Multivariate B-spline · System Identification · GPU · Cache-aware

J. Meyer

Control & Simulation Division, Faculty of Aerospace Engineering, Kluyverweg 1, 2629HS Delft, The Netherlands

E-mail: johann.meyer@gmail.com

C.C. de Visser

Control & Simulation Division, Faculty of Aerospace Engineering, Kluyverweg 1, 2629HS Delft, The Netherlands

1 Introduction

Loss-of-control in-flight (LOC-I) is the leading cause of fatalities in aviation and is a category that encompasses all accidents where the flight crew is unable to maintain control of the aircraft in-flight. This includes stalls/upsets, structural failure and adverse meteorological conditions [21]. In the period from 2009 to 2018, IATA determined that there were in total 64 LOC-I incidents in general aviation, 94% of which involved fatalities [21]. Moreover, it resulted in the highest number of fatalities compared to any other accident category, accounting for around 60% of all fatalities in general aviation [21]. In a similar study by Boeing for the same time period for commercial air transport, LOC-I was found to be the cause of around 48% of all onboard fatalities [5]. This has led to the Federal Aviation Authority (FAA) and the European Aviation Safety Agency (EASA) to mandate that flight crew are trained in handling and preventing stalls and upsets since 2019 [15] and 2016 [13], respectively. These changes have also been reflected in the best practices and guidance material by the International Air Transport Association (IATA) [20]. Furthermore, in the explanatory note by EASA, they cite LOC-I as one of the agency's highest priorities [14].

System identification plays two key roles in mitigating accidents related to LOC-I. The first is to generate accurate models of aircraft during upsets and stalls for training. Performing the training in simulators enables the simulation to be performed at a lower cost and significantly safer. The second role that system identification plays is in adaptive flight control. Adaptive flight control enables the control systems to adapt to changes in the aircraft aerodynamic model. This allows the pilot to more easily handle the aircraft despite these changes. Typically, control algorithms require knowledge of the model of the aircraft to determine the gains to be used to guarantee a certain level of performance. Nonlinear dynamic inversion (NDI) is the most used approach that is useful for nonlinear control of an aircraft. An inaccurate nonlinear model means that the nonlinear dynamics will not be sufficiently neutralised and the linear controller will not perform as desired [35]. Therefore having an accurate model at all times is critical, especially after significant changes to the model. Global smooth models are particularly important for NDI and Incremental NDI (INDI) as it ensures that there are no discontinuities in the control inputs [35] and the control effectiveness Jacobian. The discontinuities can to an extent be damped by the aircraft's inertia and actuator dynamics; however, in highly manoeuvrable aircraft, the damping may not be sufficient and may induce aeroelastic flutter. Alternatively, chattering can reduce the lifetime of components due to fatigue and wear. This last point is especially of interest for mass-produced aircraft with lower quality components, such as those found in most commercial quadrotors.

With the advent of air taxi projects like Uber Air¹ and ATTOL by Airbus², the need for adaptive control will be even more evident as these aircraft will be flying over densely-populated regions and the estimated loss of life will increase significantly. Furthermore, projects like ATTOL aim for fully autonomous operations which will require accurate models as little to no human intervention will be possible.

¹ <https://www.uber.com/us/en/elevate/uberair/>

² <https://www.airbus.com/innovation/future-technology/autonomy.html#projects>

One specific challenge for system identification is over-actuated systems with significant aerodynamic interactions. The reason this is challenging is that increasing the number of independent variables of the regression problem increases the problem size exponentially - this is typically referred to as the curse of dimensionality or, in this case, combinatorial explosion. This also increases the amount of data required to fit the model without overfitting to the data. Traditional adaptive control is implemented via gain-scheduling, which is affected by this problem to a greater extent and is not a fault-tolerant solution. In order to make the problem more manageable with hardware restraints, two approaches are common: approximating a high-dimensional space with an abstract low-dimensional space, and parallelism. These abstract spaces are in most cases arbitrary and are not a unique nonlinear mapping, limiting the interpretability of the controller. Neural networks are an example of black-boxes and are agnostic to the underlying representation, and are therefore viable options; however, nonlinear representations are often unstable [29] and, typically, there are no guarantees on performance when learning and adapting. Consequently, certification, verification, and validation are significantly harder in comparison to traditional control approaches [22].

High-performance computing and general-purpose GPU (GPGPU) computing have enabled problems to continue scaling despite power issues due to increasing clock frequencies, which was one of the primary approaches to increasing single-core performance in the beginning of computing. Brain simulation and genome sequencing have seen dramatic performance improvements due to GPGPU computing and the parallel computing paradigm, see for example the implementations in [33] for brain simulation and [28] for sequencing. Other fields related to aerospace, such as fluid dynamics and structural analysis, have also seen large improvements with GPGPU computing.

The scope of this research will be limited to function approximation of a continuous state-space where the locations of measurements are also a function of time. These choices change the structure of the optimisation problem and prevent some matrices, such as the regressor matrix, from being factorised offline. The constraints are tied to the tessellation, which is not adaptive in this case and is therefore fixed.

The aim of the research is to develop a fast algorithm for smooth online system identification that is robust to ill-conditioning by using oblique projections to simplify and improve constraint application with limited data - which is of particular concern in high-dimensional spaces. It is paramount to improve the computation speed and complexity of the constrained optimisation step to ensure the feasibility and scalability of the algorithm in online and adaptive flight control settings with limited computational power and therefore these aspects have been taken into account in this research.

The paper will first look at the current state-of-the-art in function approximators and optimisation. Thereafter some background information on multivariate B-splines, the continuity equations between the splines, and the Kaczmarz algorithm will be given. Following, a high-level overview of the Distributed Asynchronous B-spline (DAB) algorithm will be explained. Finally, the key elements of the algorithm will be expanded in the subsequent section with their complementary results.

2 Literature

This section is divided into two main parts: the first focusses on function approximators and the second on optimisation. The second part is further subdivided into two sections: unconstrained and constrained optimisation.

2.1 Function Approximators

System identification is a special subset of function approximation with the goal of modelling dynamical systems. Dynamical systems can be modelled based on first principles or they can be modelled using black-boxes, which generate an arbitrary input-output mapping that is not intended to be interpreted. Interpretability, analysability, and robustness is of key importance for certification of control systems [22]. Black-boxes also do not enable the use of a priori information when developing the model which can be used to improve accuracy and/or simplify the model complexity. The most common function approximators are fuzzy-blending, tensor-product splines, thin-plate splines, artificial neural networks, and multivariate B-splines.

First, it is necessary to distinguish between interpolation and regression. Interpolation means the function approximator must pass through all data points where as regression finds the best fit without this requirement. An implicit assumption in interpolation is that the measured data is noise-free, which is an assumption that is often wrong in experimental data and is particularly erroneous when estimating the derivative of the output. Moreover, it requires sufficient data to uniquely generate the model whereas regression can use regularisation alleviate this problem. Furthermore, in higher dimensions the amount of data required for a unique solution becomes infeasible, especially one that avoids overfitting.

Fuzzy-blending is a manual approach to smoothing and connecting splines that requires expert knowledge [35]. The smoothing must also be reperformed if the granularity of the grid is changed. Moreover, higher-order continuity cannot be achieved, which is of importance for computing the control effectiveness Jacobian. Additionally, manual approaches quickly become infeasible when the dimensionality of the problem increases.

Tensor-product splines requires taking the tensor-product of 1D polynomials. The downside of these tensor-products is that the tensor-product can create precision errors when working in higher dimensions. But most importantly, the data needs to be collected on a rectangular grid [35] and for some applications, such as flight data which is scattered, this is not possible nor desirable.

Thin-plate splines is another approach that can, in contrast to tensor-product splines, handle scattered data. The downside to thin-plate splines is that it requires a global radial basis function for each data point [35]. As a result of the global nature of the basis functions, all basis functions are required to evaluate and optimise the spline, which can quickly become computationally infeasible. This is exacerbated as the dataset increases in size and dimensions.

Artificial Neural Networks (ANN) are a set of algorithms that are modelled loosely on the neural system. These algorithms are among the most popular in recent times

due to the power and potential these algorithms have shown in domains such as reinforcement learning [30] and computer vision [27], especially when it comes to generalisation. The benefit of generalisation (extrapolation) may be of limited value for highly nonlinear aircraft and should be investigated separately. The major problems with ANNs are that they are non-convex, black-box systems, computationally dense, and require modifications to ensure that continuity and smoothness is satisfied. By computationally dense, it is meant that even if the outputs of the network are sparse, the entire network needs to be evaluated to determine the output. Spiking Neural Networks are a new generation of neural networks for neuromorphic hardware that will be able to exploit this sparsity. Moreover, since they are black-box systems, the entire ANN needs to be evaluated multiple times to estimate derivative information which can be computationally expensive depending on the size of the network. Furthermore, small changes in the input do not necessarily mean small changes in the output as a result of the nonlinearities. It is also much harder to ensure all constraints are satisfied throughout the entire domain as they suffer from problems such as catastrophic forgetting (see [29] for more details). Catastrophic forgetting is also more apparent for online learning as the data is highly correlated. For example in reinforcement learning, nonlinear function approximators, like neural networks, are unstable as a result of this and have required modifications like memory buffers of past data and copies of the networks to stabilise the algorithms [30]. The non-convexity of ANNs lead to slower training and gradient-based methods cannot guarantee global optimality; however, first-order algorithms and local minima are often sufficient in practice. The size of the networks and the plethora of hyper-parameters that are associated with them are empirically determined and provide no guarantees of their performance in-situ and complicates the design and certification process [22].

Multivariate Simplex and Simplotope B-Splines are linear-in-the-parameters function approximators that can be used to model nonlinear data. They can be solved using Linear Least Squares and can therefore be optimised using convex solvers. Since it is a linear-in-the-parameters function approximator, there are more guarantees for the stability of the approximator and the mathematical tools for analysis of the approximator are more mature. The curse-of-dimensionality is present since no nonlinear map is used. It is, nevertheless, possible to use dimensionality reduction techniques, such as autoencoders [17], principle component analysis, tensor networks [24], and hashing functions to alleviate this but, as with any compression technique, some data loss will occur, accuracy will be decreased, and sparsity will be reduced. The extent to which this occurs can be managed with appropriate trade-offs. An alternative approach to reducing the problem size, is using null-space projection [9] of the constraint matrix; however, this approach still reduces the sparsity. Nonetheless, the interpretability is lost, if a dimensionality reduction technique is used. Moreover, it should remain static to avoid invalidating the function approximator that relies on that input representation. The curse-of-dimensionality is additionally abated by the use of local basis functions, as these result in a very sparse problem structure that can be efficiently evaluated. The optimisation of the parameters can also occur in a distributed and parallel fashion [1]. Continuity and differential constraints can be applied as shown by De Visser [35] and a priori knowledge can be incorporated using physical constraints as shown by van der Peijl [34] and Huisman [19].

2.2 Optimisation

Optimisation is a large field and a wide body of literature exists. The optimisation problem for function approximation is typically limited to the Least Absolute Error (LAE), Least Squared Error (LSE), and Huber Loss. LSE is the most popular for regression and quadratically penalises the error. LSE has a smooth cost function and has an analytical solution. LAE is useful for sparse solutions but requires the use of linear iterative solvers that can handle non-smooth cost functions [3]. An additional benefit of LAE is that outliers are not as heavily penalised as for LSE and is therefore more robust. Huber loss is a hybrid of LSE and LAE and aims to improve the robustness of LSE and the convergence rate of LAE; however, it requires switching between two different optimisers.

2.2.1 Unconstrained Optimisation

Unconstrained optimisation can be split into two main methods: gradient-based and derivative-free methods. Derivative-free methods are useful when the loss function is noisy or when evaluating the gradient is expensive or impossible. Additionally, they are useful for non-convex optimisation as they can escape local minima. Gradient-based methods are typically preferable as the gradient information can be used to direct the optimisation and lead to faster convergence, especially in higher dimensions where the curse of dimensionality can make it infeasible to effectively explore the domain. Moreover, the gradient information can be used to determine the optimality of the solution. Gradient-based methods convergence rates are sensitive to the scaling of the loss function and as a result matrix conditioning is important. gradient descent and coordinate descent are the two foundational algorithms when looking at gradient-based methods. Most algorithms are variations of these two where the changes are the step size and the direction of descent. Popular gradient-based algorithms are steepest descent, Newton's method, quasi-Newton methods, conjugate gradient descent, and momentum-based gradient descent. Momentum-based methods - such as ADAM [26], Nesterov accelerated gradient [31], RMSprop, and AdaGrad [10] - are popular in machine learning and increase the rate of convergence significantly but may reduce generalisation accuracy [25]. Newton's method results in one-step convergence but requires the inversion of a matrix, which can be slow and numerically inaccurate as well as being badly conditioned in higher dimensions.

2.2.2 Constrained Optimisation

Constrained optimisation are often meta-algorithms in the sense that they require unconstrained optimisation algorithms to be executed during some of the steps, which direct the unconstrained optimisation to a constrained optimal solution.

Kaczmarz method is a special case of Projection onto Convex Sets (POCS), where the method of alternating projections is applied to affine constraints. In the field of image reconstruction, it is referred to as the algebraic reconstruction technique. Dykstra's method is a generalisation of Kaczmarz for convex sets [11] and translated

convex sets [12] and guarantees that the algorithm converges to the orthogonal projection within the constraint set. A parallel version of Dykstra's algorithm is given by Gaffke and Mathar [16]. Kaczmarz method can also be used as an unconstrained optimisation method for solving the least squares problem; however, it excels when the problem is sparse as in the constraints between splines. It should be noted that when used for constrained optimisation, the algorithm does not take into account the loss function and does not result in the constrained optimal solution. It does; however, introduce some flexibility into the algorithm by enabling the algorithm to be split into two orthogonal steps: cost optimisation and constraint satisfaction.

Null-space projection uses a projection to transform a constrained problem into an unconstrained problem. The null-space projection can be computed offline; however, it is susceptible to numerical precision errors. Additionally, the Hessian of the unconstrained problem is no longer block-diagonal, which significantly increases the cost of solving the unconstrained problem. A numerically robust online form of null-space projection is presented by Zhu [39], where the constraints are applied by differing initial conditions.

Lagrange multiplier methods are a family of algorithms that introduces the concept of dual variables and the dual function. The duality gap is the difference between the primal and dual function and can be used to draw conclusions about the optimality of the constrained solution [6]. The most intuitive version is the penalty method that involves adding a penalty term for violated constraints to the loss function. The penalty multiplier is iteratively increased until it converges to the constrained optimal solution. This approach can sometimes lead the cost function to focus on constraint satisfaction over optimising the original cost function. For quadratic problems the Lagrange multiplier method has an analytical solution; however, computationally this method is poorly conditioned and slow due to the large matrix inversion. An implicit iterative method by [2] is often faster and does not explicitly store the Lagrange multipliers. Dual ascent, dual decomposition, and the method of multipliers alternate between optimising the primal and dual functions. Dual decomposition is the parallelised version of dual ascent. Dual ascent is more sparse than the method of multipliers but requires a smooth and differentiable cost function. The method of multipliers introduces an additional hyper-parameter - the penalty factor - and it enables non-differentiable cost functions with possibly infinite values to be used [6]. The alternating direction method of multipliers (ADMM) makes the problem separable by adding auxiliary constraints and an additional optimisation step [6]. Additional optimisation steps can only be added if the steps are orthogonal [7]; however, this also means these steps can be run in parallel. Wei proposed a distributed [37] and an asynchronous [38] version of ADMM; however, these have yet to be applied to multivariate B-splines.

2.3 Optimisation in the Context of Multivariate B-splines

The first use of multivariate B-splines for aerodynamic modelling was by De Visser [35] and he used the recursive least squares formulation by Zhu [39]. Dual ascent has been used for multivariate simplex B-splines by De Visser in a sequential [9] and distributed manner [8]; however, that was for wavefront reconstruction, which

involves continuous measurements at fixed locations and enables a different set of assumptions to be made on the structure of the problem. A distributed version of ADMM has been used for multivariate simplotope B-splines by van den Aarssen [1] and Hooij [18] for aerodynamic model identification. Hooij expanded on van den Aarssen's work by introducing a distance weighting function for using pre-failure data in online model identification. De Visser [8] and van den Aarssen [1] both utilise a hybrid approach where the grid is split into partitions and null-space projection is used internally in a partition and another approach is used for inter-partition constraints. The disadvantage of null-space projections is that they result in matrix fill-in and are numerically unstable for high dimensions and large partition sizes. Awanou [2] used the implicit Lagrangian iterative solver for solving partial differential equations using multivariate B-splines. Karagöz's [24] tensor-network B-splines are equivalent to simplotope B-splines by Van Den Aarssen [1] and multiplex B-splines by Visser [36]. Karagöz's [24] approach to high-dimensional function approximation is a form of compression of the network to reduce its memory footprint at the cost of accuracy; however, at the same time the approximator is no longer interpretable and continuity is not taken into account. This research aims to improve the robustness of constrained online optimisation while ensuring that the algorithm can be effectively executed in an online setting.

3 Background

The background is divided into three main sections: theory on multivariate orthotope B-splines, Kaczmarz's algorithm, and caches.

3.1 Multivariate Orthotope B-splines

A simplex t_j is a geometrical structure that minimally spans a set of n dimensions and is composed of exactly $n + 1$ vertices where the edges are given by simplices of one dimension lower. Simplices can be connected to each other at their edges. A collection of these connected simplices is called a tessellation \mathcal{T} .

Each simplex has a local coordinate system known as the barycentric coordinate system. A useful property of barycentric coordinates is that points that lie outside the simplex have one or more negative barycentric coordinates. To convert Cartesian coordinates to barycentric coordinates, the transformation matrix $T \in \mathbb{R}^{n \times n}$ for a simplex should be formed. The barycentric coordinates $b \in \mathbb{R}^{n+1}$ can then be computed using Eq. 2 and Eq. 3. The vertices v of the simplex are numbered using the superscript and their Cartesian coordinates are indexed by their subscript.

$$T = \begin{bmatrix} v_0^1 - v_0^0 & \cdots & v_0^n - v_0^0 \\ \vdots & & \vdots \\ v_{n-1}^1 - v_{n-1}^0 & \cdots & v_{n-1}^n - v_{n-1}^0 \end{bmatrix} \quad (1)$$




Fig. 1 Simplotopes can be described as N-simplices and N-orthotopes at the extremes. The figure is a 2D projection of the N-dimensional simplotopes. N-simplices are generated by adding a vertex in the n^{th} dimension. N-orthotopes are created by duplicating the (N-1)-orthotope in the n^{th} dimension and connecting the respective vertices or, alternatively, as a tensor-product of 1D-simplices. Adapted from [1].

$$\begin{bmatrix} b_1 \\ b_2 \\ \vdots \\ b_n \end{bmatrix} = T^{-1} \begin{bmatrix} x_0 - v_0^0 \\ x_1 - v_1^0 \\ \vdots \\ x_{n-1} - v_{n-1}^0 \end{bmatrix} \quad (2)$$

$$b_0 = 1 - b_1 - b_2 - \dots - b_n \quad (3)$$

The Bernstein basis functions are given by the multinomial theorem (Theorem 1). It is a stable basis that sums to 1. Additionally, it is local because the basis functions are zero outside of the simplex. Each basis function B_κ^d is scaled by its B-coefficient c_κ to approximate any polynomial, as given by De Boor's theorem (Theorem 2). $\mathbf{B}^d \in \mathbb{R}^{l \times \hat{d}}$ is the matrix formulation of the summation in Eq. 7, where l is given by the number of barycentric coordinates being used and \hat{d} is given by Eq. 4. In the case of the theorem, $l = 1$. $\vec{1}$ is a column vector of ones. The B-coefficients each belong to a specific spline t_j and may be distinguished by a superscript to indicate this membership. The barycentric coordinate of each B-coefficient with respect to a spline of degree d is given by Eq. 10. $\kappa \in \mathbb{N}_0^{n+1}$ is a multi-index. An additional notation used for multi-indices is \mathbb{M}_d^{n+1} , where the superscript denotes the set of \mathbb{N}_0^{n+1} subject to the constraint that the norm of the multi-index is equal to the subscript d and the number of elements in the set is determined by filling in the subscript in Eq. 4.

$$\hat{d} = \frac{(d+n)!}{d!n!} \quad (4)$$

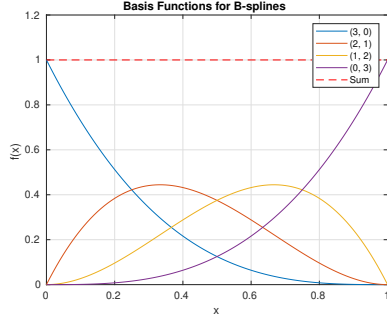


Fig. 2 Bernstein basis functions for a 3rd-order 1D spline.

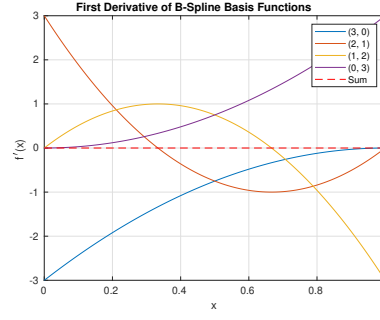


Fig. 3 1st-order derivative of the Bernstein basis functions for a 3rd-order 1D spline.

Theorem 1 (Multinomial Theorem for Barycentric Coordinates)

$$(b_0 + b_1 + \dots + b_n)^d = \sum_{|\kappa|=d} \frac{d!}{\kappa_0! \kappa_1! \dots \kappa_n!} b_0^{\kappa_0} b_1^{\kappa_1} \dots b_n^{\kappa_n} \quad (5)$$

$$= \sum_{|\kappa|=d} \frac{d!}{\kappa!} b^\kappa \quad (6)$$

$$= \sum_{|\kappa|=d} B_\kappa^d(b) \quad (7)$$

$$= \mathbf{B}^d \vec{1} \quad (8)$$

$$= 1 \quad (9)$$

Theorem 2 (De Boor's Theorem) Any polynomial $p(x)$ of degree d can be written as a sum of basis functions B_κ^d scaled by its B -coefficient $c_\kappa^{t_j}$:

$$p(x) = \sum_{|\kappa|=d} c_\kappa^{t_j} B_\kappa^d(b(x))$$

$$b(c_\kappa) = \frac{\kappa}{d} \quad (10)$$

Since the basis functions are local to a spline, the evaluation of the multivariate B-spline for any point only requires one to evaluate the function approximator for the spline that contains the point. During regression, updates only need to be applied locally. For second order methods, this manifests itself as a sparse block-diagonal structure which significantly reduces the computational complexity, especially for large tessellations.

Simplotope B-splines - also known as multiplex B-splines in [36] - is a modification to simplex B-splines where the tensor-product of lower-dimensional orthogonal simplices are used instead of one high-dimensional simplex. The more layers in a simplotope spline, the fewer the number of constraints for the same grid size; however, the number of parameters per simplotope also increases with the number of layers. More parameters per simplotope means larger block sizes in the block diagonal data matrix.

The benefit of more layers is that the total number of coefficients are reduced, as fewer coefficients are duplicated and there are fewer simplotopes in the same triangulation and therefore fewer constraints to be applied.

Simplotope equations similar to their simplex counterparts and are denoted with π to distinguish them. The multi-index κ is similar to the those for simplices; however, instead of one multi-index, there is one for each simplex layer l , which are concatenated to form κ (i.e. $\kappa = [\kappa_1, \dots, \kappa_l]$). The same applies to the barycentric coordinates b . Eq. 11 is used to evaluate a simplotope. B_{κ_i} in Eq. 12 is the B-form for the simplex in layer i with the subset of the multi-index κ for that layer.

$$\pi(b) = \sum_{\substack{|\kappa_i|=d_i \\ \forall i \in [1, l]}} c_{\kappa} B_{\kappa}(b) \quad (11)$$

$$B_{\kappa}(b) = \prod_{i=1}^l B_{\kappa_i}^{d_i}(b_i) \quad (12)$$

Orthotope B-splines is a special case of simplotope B-splines, where the lower-dimensional simplices are all one dimensional (see Fig. 1). Orthotopes are also referred to as hyperrectangles in literature. Orthotope B-splines are beneficial as they simplify computation of neighbouring splines, high-dimensional tessellations, spline membership, and indexing. Additionally, orthotopes have fewer neighbours than their equivalent simplex form and, in higher dimensions, simplices become *sliver simplices* - as described by De Visser [35] - which are difficult to fill with data. Moreover, the approximation power of 1-dimensional simplices is higher than that of higher-dimensional simplices as couplings are created between the dimensions when applying continuity. The use of 1-dimensional splines also enables the tensor-product to be evaluated in parallel.

See De Visser [35] for more detailed information on multivariate simplex B-splines and Van den Aarssen [1] for more details on simplotope B-splines.

3.1.1 Constraints

There are two main types of constraints that are applicable to aerodynamic function approximation. The first being continuity constraints between the simplices in the tessellation and the second being a priori knowledge of the problem.

Continuity constraints are particularly important for control systems that are based on the estimated aerodynamic model as discontinuities or also lack of differentiability between simplices will either limit the control algorithms that can be applied or limit the performance of the control system. The discontinuities and lack of differentiability is also an artifact of using splines and not necessarily related to the underlying function that is being approximated. $v_*^{t_i}$ is the barycentric coordinate of the out-of-edge vertex for simplex t_i with respect to simplex t_j . The non-zero multi-index for the out-of-edge vertex for simplex t_i is denoted as $\kappa_*^{t_i}$. Finally, γ is a multi-index. A visual depiction of the coefficients involved in the constraints for the 1D case is given in Fig. 19. It should be noted that the continuity matrix is purely a function of the tessellation and its connectivity and is not dependent on the data being regressed. This means that this

matrix is always constant. A consequence of applying the spline continuity constraints is that the local model becomes a global model. The distance the constraints propagate through the model is related to the number of degrees of freedom of each simplex [35]. Therefore, higher-order models with lower-order continuity leads to more local models. The upper-bound on the number of constraints is given by Eq. 13 multiplied by the number of shared edges. This is also the number of constraints generated by following the algorithm in Theorem 3 and, as a result, some constraints are redundant which has repercussions for some solvers. If the nullity of the constraint matrix is zero, only the trivial solution exists.

Theorem 3 (Continuity Equations)

$$c_{\kappa_0, \dots, \kappa_n}^{t_i} = \sum_{|\gamma|=m} c_{(\kappa_0, \dots, \kappa_n) + \gamma}^{t_j} B_{\gamma}^m(v_*^{t_i}), \quad \{\forall m \in \mathbb{N}_0 | 0 \leq m \leq r\}$$

subject to:

$$\begin{aligned} \kappa_0^{t_i} + \dots + \kappa_n^{t_i} &= d \\ \kappa_*^{t_i} &= m \\ \kappa_*^{t_j} &= 0 \end{aligned}$$

$$\hat{r} = \sum_{m=0}^r \frac{(d-m+n-1)!}{(d-m)!(n-1)!} \quad (13)$$

The application of constraints for simplotopes is very similar to that of simplex B-splines; however, requires a slight modification. The process of generating the continuity constraint as in Theorem 3 is only applied to the simplex that contains the out-of-edge vertex. The tensor-product is then applied to all permutations of the simplices in the other layers. The number of continuity constraints for an edge of a simplotope is given by Eq. 14, where \hat{r}_* is given by Eq. 13. \hat{r}_* is the number of constraints for the out-of-edge simplex. \hat{d}_i and \hat{d}_* are the number of parameters in layer i and the out-of-edge simplex, respectively. The propagation of constraints in orthotope B-splines can be more easily controlled than in simplex B-splines due to the more regular structure.

$$\hat{r}_{\pi} = \hat{r}_* \frac{\prod_{i=1}^l \hat{d}_i}{\hat{d}_*} \quad (14)$$

A priori knowledge can for example be incorporated into the function approximation problem by including constraints due to physical constraints associated with the aircraft. However, the simplex B-splines operate using barycentric coordinates and, in order to be able to apply these constraints, the physical constraints in Cartesian coordinates need to be transformed to barycentric constraints. This approach to constraint application is referred to as physical splines in literature [19] and [34]. It is also possible to apply constraints based on the directional derivatives as in [9] using the De Casteljau matrix in Appendix A.

3.2 Kaczmarz's Method

Kaczmarz's method [23] optimises the cost functions given by Eq. 15 and Eq. 16 when A is overdetermined and underdetermined, respectively. θ_0 are the initial model parameters and y is the vector of target values. The iterative equation for orthogonal projections is given by Eq. 17. A_j is the j^{th} row of A . j can be selected in any order but is commonly chosen either stochastically or cyclically. Note γ is an optional relaxation parameter.

The method can be viewed from two perspectives. The first perspective is that it involves iteratively projecting a vector onto a linear constraint set - i.e. it is a special case of projection onto convex sets (POCS). The second perspective is that of gradient descent, where the numerator is the gradient of the cost function and the optimal step-size for the i^{th} iteration α_i^* is given by Eq. 18. The algorithm, in general, does not converge in n steps, where n is the number of rows in the matrix A . This is because the rows are not guaranteed to be orthogonal nor are the search directions A -orthogonal, as is the case in the conjugate gradient method[32]. The algorithm is computationally beneficial when the rows of A are sparse making it cheaper and more precise than a line search. Additionally, the solution using Kaczmarz's method for an underdetermined system results in a higher precision than other iterative approaches. This is of importance, for example, when applying equality constraints.

$$\underset{\theta}{\text{minimise}} \quad \|y - A\theta\|_2^2 \quad (15)$$

$$\begin{aligned} &\underset{\theta}{\text{minimise}} \quad \|\theta - \theta_0\|_2^2 \\ &\text{subject to } A\theta = y \end{aligned} \quad (16)$$

$$\theta_i = \theta_{i-1} + \gamma \frac{(y_i - A_j \theta_{i-1}) A_j^T}{\|A_j\|_2^2} \quad (17)$$

$$\alpha_i^* = \frac{1}{\|A_j\|_2^2} \quad (18)$$

3.3 Caches

The fact that processor performance improves much faster with time than memory performance is known as the processor-memory performance gap. Memory hierarchies is one approach to mitigate the loss of performance due to this gap and works by organising system memory into a hierarchy of faster but progressively smaller and more expensive memories. Data that is accessed more frequently or more recently is placed in the faster levels to reduce latency. The data from slower main memory is transferred to the faster caches in blocks of contiguous data known as cache lines to reduce routing complexity. If the data is not reused, these blocks will be evicted or, if the data is not contiguous, cache pollution may occur - reducing the effective size of the cache. Cache thrashing occurs when the cache is too small to hold the required data leading to data being evicted before it can be reused.




Fig. 4 A simplified cache system is illustrated, where the # denotes a location in main memory. The red squares indicate the memory addresses that will be modified by the core. Whenever an element in a cache line is written to, the other caches are informed by the write-invalidation policy that it must re-fetch its cache line from main memory or the lowest shared cache level. False sharing is the case where only the green elements are written to by a core. Even though no relevant data has changed for the core, the data must be re-fetched leading to a large latency.

In multicore systems with multiple caches, cache coherency needs to be ensured. Cache coherency is the process of ensuring all caches agree on the same value for a memory address. First, it is necessary to ensure that all data in a cache is eventually written to main memory. The two main approaches to ensuring this is the write-through and write-back protocols. Write-through is simpler and makes all writes to a memory address in the cache immediately to the copy in main memory. This approach is typically found in GPUs and has the disadvantage of using a lot of possibly unnecessary memory bandwidth, if many writes occur. Write-back is more complex and typically found in CPUs. It involves keeping track of whether the data has been modified and defers writing to main memory until the cache block is evicted or requested by another core. If another core requires the data, this approach leads to a longer latency as data must first be written to main memory before it can be requested by the other cache. One common approach to cache coherency is using write-invalidation. Write-invalidation involves broadcasting to all cache copies that the data that they have for a cache line is no longer valid and should be reloaded from main memory. This leads to another problem, which is known as false sharing. False sharing is when different cores are using different data in the same cache lines (see Fig. 4). When either core writes to the data, the other core needs to reload the data from main memory, even though its data has not been modified.

Scratchpad memory is a form of fast memory that is used to store intermediate calculations that do not need to be stored in main memory. It is typically - but not necessarily - explicitly managed by a programmer. It is often similar to L1 cache; however, it does not require the overhead of cache coherency protocols or writing to main memory as the data is local to a core. This reduces contention for the memory bandwidth. In Nvidia GPUs, the scratchpad memory is known as shared memory and is sometimes shared with the L1 cache and for the Turing and Ampere architectures it can be set to maximum size of 64KB and 100KB, respectively. The scratchpad memory is local to a core for CPUs or Streaming Multiprocessor (SM) for GPUs (see Fig. 5) and cannot be used for inter-core communication.

⁴ Intra-warp communication is available using warp-level primitives.




Fig. 5 Each rectangle indicates a physical separation between resources. In red is RAM, in blue is cache and green is used for computing cores. CPUs are based on a smaller number of large complex cores. In contrast, GPUs are based on a large number of simpler cores that are grouped into Streaming Multiprocessors (SM) - in Nvidia terminology - that share certain resources, like an L1 cache. Each of the tiny cores in an SM can only communicate with other cores within its SM using shared memory or with all cores via main memory.⁴

Sorting the data in memory according to expected access patterns leads to the expected latency being closer to the cache latency rather than the main memory latency, which can differ by several orders of magnitude. Additionally, avoiding false sharing is critical to avoid unnecessary cache invalidation. Finally, effective use of scratchpad memory can reduce the load on the memory bandwidth allowing relevant data to be retrieved with a lower latency.

4 DAB Algorithm

In this section an overview of the key features of the Distributed Asynchronous B-spline Algorithm (DAB) will be given. A graphical overview of the algorithm is given in Fig. 6. First, a maximum allowable continuity is set to limit continuity propagation. Thereafter the reasoning behind oblique projections is explained. Finally, the need for a specific memory layout is given.




Fig. 6 An overview of the DAB algorithm. The algorithm takes as input an unconstrained solution and distributes the data of each vertex to a different processor. The constraints are then applied to each vertex through oblique projections. The output of the algorithm is a constrained solution.

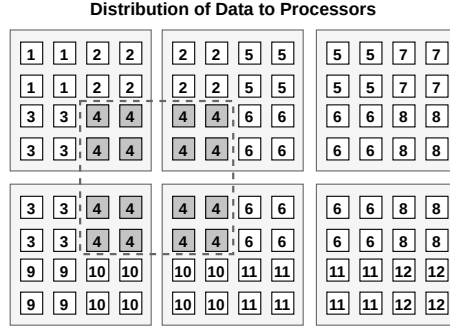


Fig. 7 Each vertex (indicated by dashed line) can be distributed to a different processor (indicated by number). The example is of a third-order B-spline in two dimensions. In higher dimensions, the blocks become orthotopes.

The maximum continuity of a spline of degree d is given by Eq. 19. Graphically, this means the B-coefficients involved in the continuity cannot pass the halfway mark with respect to the spline for that layer or the continuity will propagate throughout the entire tessellation. This propagation has two negative impacts on the algorithm. The first is that the continuity will take time to propagate from one side of the tessellation to the other. This can be abated by multi-grid methods (see for example[4]); however, it is computationally better to keep the modifications in the model local. Additionally, this is physically more intuitive and justifiable. Moreover, the parallelisability of the algorithm is impacted as the algorithm must either utilise red-black ordering or include auxiliary constraints to ensure correctness. While not necessary, the algorithm simplifies significantly if the degree of the model is set based on the continuity required using Eq. 19 as edge cases are eliminated (see the boundaries of Fig. 7). The edge cases can also be alleviated by padding the domain with splines with their weights set to 0. This may, however, lead to slower convergence than handling the edge cases directly. The algorithm is asynchronous in that each vertex requires no communication with any other vertex, as shown in Fig. 7. This is beneficial in the offline/batch setting and when large high-dimensional tessellations are being used. Within each vertex, the algorithm is parallel but requires all the projections for a direction to complete before changing direction.

$$\mathbb{C}^{\max} = \left\lfloor \frac{d+1}{2} \right\rfloor - 1 \quad (19)$$

The B-spline cost function is given by Eq. 20. Gradient descent with ADAM [26] with early stopping is used to optimise the unconstrained problem for fast convergence with low data coverage, which can impact the stability of second-order algorithms or otherwise lead to overfitting. Kaczmarz's method is used to apply the constraints and is solving Eq. 16, with θ_0 given by the unconstrained optimal solution. Kaczmarz's method gives a solution that satisfies all constraints but does not provide the optimal constrained solution as it does not take into account the regressor matrix. Kaczmarz's method is equivalent to dual ascent with the second term of the derivative being changed from the Hessian $B^T B$ to the identity matrix (compare Eq. 21 and Eq. 22).

This means the cost achieved by Kaczmarz's method is not cost-optimal based on the available data B . It is, however, possible to reduce the cost by using a few steps of dual ascent prior to applying Kaczmarz's method. With more dual ascent iterations, the solution will approach the optimal constrained solution. Lagrange multipliers can be interpreted as "forces" applied to the B-coefficients. When the Lagrange multipliers are large, it is indicative of overfitting the model to the obtained data or that the desired continuity level is not appropriate for the current model. If the data is not the issue, then the latter issue can be improved by reducing the size of the orthotopes or by reducing the desired order of continuity. Moreover, dual ascent is a global approach as the basis functions (see Fig. 2) are non-zero away from vertices and need to take into account the data from the entire orthotope with either 2 or more dimensions, or first-order continuity. The constrained solution obtained by Kaczmarz's method is the one with the minimum "forces" applied and results in less oscillatory behaviour at apparent discontinuities between splines. The oscillatory behaviour is akin to the Gibb's phenomenon. The discontinuities are termed as apparent as overfitting to limited data can result in pseudo discontinuities, which may not be representative of the true underlying model. These discontinuities are indistinguishable from actual discontinuities until sufficient data is available and therefore modelling them may or may not be desirable.

$$\begin{aligned} & \underset{\theta}{\text{minimise}} \quad \|y - B\theta\|_2^2 \\ & \text{subject to} \quad H\theta = 0 \end{aligned} \quad (20)$$

$$\frac{\partial f}{\partial \theta} = (B^T B)^{-1} B^T y - IH^T \lambda \quad (21)$$

$$\frac{\partial f}{\partial \theta} = (B^T B)^{-1} B^T y - (B^T B)^{-1} H^T \lambda \quad (22)$$

Oblique projections are used to extrapolate or weight the estimates of the B-coefficients such that poorly-estimated coefficients do not influence the well-estimated coefficients when applying continuity. Fig. 8 illustrates this by showing how the right spline pulls the left splines coefficients down. Fig. 9 illustrates how oblique projections avoids this and achieves better accuracy on the training dataset and, in principle, the validation dataset.

The optimisation algorithm applies the constraints by iterating through the out-of-edge dimensions sequentially. Memory on parallel computing systems divide the memory into multiple memory chips to enable parallel memory access; however, when the out-of-edge dimension changes, the naive ordering (Fig. 10) results in all parallel threads reading from the same memory chip and requires each access to be sequentially handled. By padding the memory, each thread can be guaranteed parallel access to its data.

Additionally, memory on computing systems are organised in a hierarchical fashion from fastest and most expensive to slowest and cheapest - where the difference can be several orders of magnitude. To make efficient use of the hierarchy, data is often transferred in contiguous batches to the faster memory. If all memory accesses are closer together, better use of the memory hierarchy is made. Sorting the B-coefficients




Fig. 8 The function has been regressed for $x < 0.5$ on the reference function. For orthogonal projections, it can be seen that the predictions from the left orthotope are pulled down to match the right orthotope, even though no data has been acquired for $x > 0.5$.




Fig. 9 The function has been regressed for $x < 0.5$ on the reference function. For oblique projections, it can be seen that the predictions from the right orthotope are pulled up to match the left orthotope, thereby improving the accuracy on the reference function.




Fig. 10 Two methods to map the B-coefficients of an orthotope to memory banks. Consecutive memory addresses are mapped to consecutive memory banks. The naive ordering results in a 4-way bank-conflict when the vertical direction is the out-of-edge direction. The conflict-free ordering pads the data and therefore uses more memory but avoids the bank conflict across the D2 line (in red) leading to parallel accesses to the B-coefficients and faster overall execution.

in memory according to ones that are frequently used together for continuity and regression can ensure that all data in the caches are relevant. In high-level programming languages, less control is available; however, rows or columns of matrices are still typically stored contiguously to minimise overhead and the cache repercussions of sharing cache lines is still present.

5 Oblique Projections

This section builds upon the view of Kaczmarz's method as an iterative projection algorithm, as mentioned in Section 3.2. The original algorithm typically assumes orthogonal projections; however, this approach is not necessarily optimal as the estimate of the parameters in a neighbouring spline may be better or worse depending on the

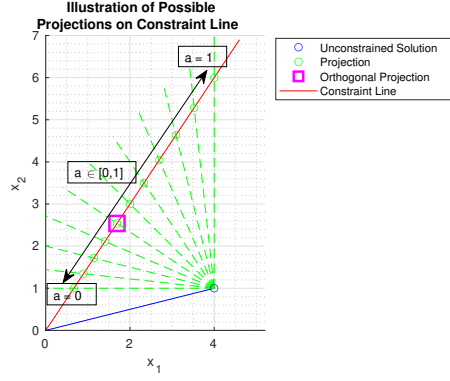


Fig. 11 The projection direction can be modified by varying the parameter a in Eq. 23, which changes the gradient v . Horizontal or vertical projection prevents a change in x_2 and x_1 , respectively. This is useful to improve the robustness of the algorithm in the case that the uncertainty of one parameter is larger than another. A projection direction with $a = 0.5$ is not necessarily equivalent to the orthogonal projection direction.

data coverage. The vector v in Eq. 23 can be used to control the projection direction. A constraint is always between two splines and the relation in Eq. 23 always holds; however, in order to generalise the relation, the vectors need to be extended to the size of constraint equation and the elements are one in the left vector if the coefficients belong to the left spline and similarly for the right vector. a is then the normalised weight for the spline. The m^{th} orthogonal projection direction H_m is given by the continuity relation for m^{th} -order continuity. The rows of H are identical for the same continuity order but differ only in the indexing function. As a result, it was decided to subscript H with m instead of i (see the indexing function in Appendix C for more details). For clarity, the splines are referred to as the left and right spline with respect to the edge; however, this is a general relation for all dimensions. The projection for the m^{th} continuity order is denoted as \mathbb{P}^m and is computed using Eq. 25 and Eq. 24. $\theta_k^{i_l^m, i_r^{0:m}}$ denotes the value of the model parameters at iteration k at indices i_l^m and i_r^0 through to i_r^m , in that order. The indices i use a subscript to differentiate between the parameters of the two splines and the superscript determines the offset from the edge. The operator \odot indicates element-wise multiplication. The parameters of the model are updated using Eq. 26.

$$v = a \begin{bmatrix} 1 \\ 0 \end{bmatrix} - (1 - a) \begin{bmatrix} 0 \\ 1 \end{bmatrix} \quad (23)$$

$$w^m = \left[w_l \overbrace{w_r \cdots w_r}^{m+1} \right]^T \quad (24)$$

$$\mathbb{P}^m = \frac{H_m \theta_k^{i_l^m, i_r^{0:m}}}{H_m H_m^T} H_m^T \odot w^m \quad (25)$$

$$\theta_{k+1}^{i_l^m, i_r^{0:m}} = \theta_k^{i_l^m, i_r^{0:m}} - \frac{\mathbb{P}^m}{w_l + w_r} \quad (26)$$




Fig. 12 This figure illustrates poor data coverage, as some splines have no data.




Fig. 13 This figure illustrates better data coverage, as all splines have some data.




Fig. 14 This figure is generated - using 5-fold cross validation on the data presented in Fig. 13 - to illustrate how overfitting of the unconstrained solution can lead to worse validation accuracy, which cannot be solved by higher-order continuity. Continuity order has little impact on validation accuracy for the proposed algorithm.

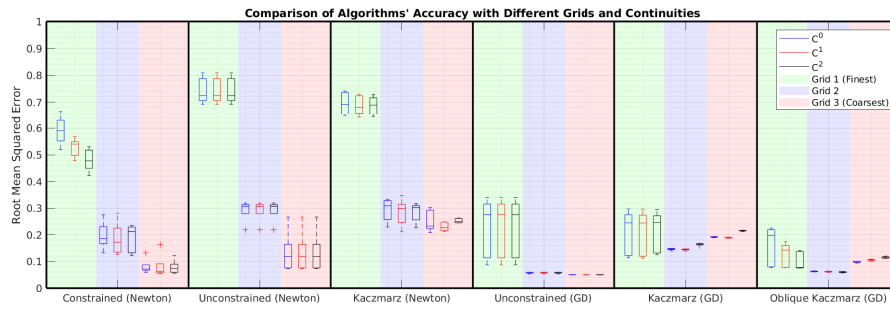


Fig. 15 This figure is generated - using 5-fold cross validation on the data presented in Fig. 12 - to illustrate how oblique projections can be used to improve validation accuracy while reducing the variance in the estimated coefficients when data coverage is low in some regions of the domain. Coarser grids are better and continuity has little impact on validation accuracy for the tested dataset. The variance in the accuracy of the model for Kaczmarz's method is generally lower and is therefore more reliable for online use.

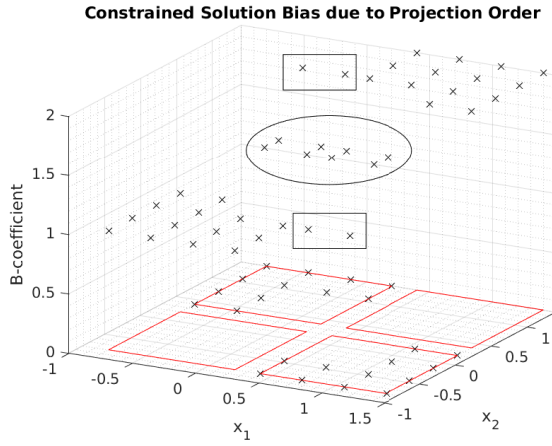


Fig. 16 This figure illustrates how the order of the projection directions can influence the constrained solution. The first projection direction is x_1 . The coefficients in the rectangles illustrate that the first projection extrapolates the splines in these directions first. Projection in the x_2 direction then averages the values in the ellipse. The bias is apparent from the rectangular shape in the ellipse. If the projection order is swapped, the ellipse will be rotated by 90 degrees. It is interesting to note that oblique projections enable continuity to be propagated in the diagonal direction, which would not be possible if the constraints to neighbouring splines were to be omitted due to insufficient data.

In order for the weighting of the oblique projections to operate as intended, the projections must converge in a dimension before it can proceed to the next dimension. Orthogonal constraints, as in Section 6, provide a clear benefit as only one iteration per dimension is required. The weightings do not change and are determined by the quality of the data in the spline or by whether or not the spline contains valid or sufficient data.

6 Orthogonal Constraints

Orthogonal constraints have a geometrical interpretation in the case of spline continuity but it is first necessary to graphically review the spline continuity formulation in Theorem 3. In [35], the continuity equations are formulated asymmetrically across the spline boundary, as can be seen in Fig. 19. This asymmetry means that once \mathbb{C}^1 has been applied, \mathbb{C}^0 is no longer true. This case will, hereon, be referred to as constraint divergence. In contrast, the orthogonal formulation - as in Fig. 20 - updates both splines' parameters thereby ensuring that lower-order continuity is still maintained. For even and odd continuity-orders, the coefficients are anti-symmetric and symmetric about the spline boundary, respectively. Although not apparent from Fig. 19 and Fig. 20, the computation and number of memory accesses approximately doubles asymptotically, in the 1D case (compare Fig. 21 and Fig. 22). In the higher-dimensional case, the number of operations required for orthogonality results in a significant overhead - as the orthogonality means that all parameters for all connected splines around a vertex need to be updated to avoid constraint divergence. Using the 1D orthogonal equa-




Fig. 17 The ratio of the number of multiplications required per iteration when comparing the original 1D formulation to the orthogonal 1D formulation. The ratio never exceeds two and the original version requires typically more than 1 iteration due to the non-orthogonality of the constraints. The ratio can also be deduced by comparing the number of non-zero values in Fig. 21 and Fig. 22. The jaggedness of the line arises due to the two zero values in the centre of Fig. 22 at every even continuity order. Oblique projections require convergence in the 1D case before changing direction. It is therefore never beneficial to use the original constraints in 1D.




Fig. 18 The required number of iterations for convergence for the original formulation is determined by the distance of the unconstrained solution to the constrained solution, which is not known a priori. However, it is clear from the figure that the orthogonal constraints are always computationally cheaper than two iterations of the original formulation (see also Fig. 17) and the slope of the line increases with the continuity order.

tions instead of the original equations results in the algorithm converging in one pass through per dimension. The original version would need to converge in one iteration to perform better from an arithmetic complexity perspective (see Fig. 17), which is guaranteed to not occur due to non-orthogonality. Fig. 17 illustrates the ratio of the number of multiplications required per iteration for the original to the orthogonal case. The right term is for the original formulation which converges in t iterations with the per iteration complexity given by Eq. 27. The orthogonal version requires only one iteration and an upper-bound on its complexity is given by Eq. 28. The bound is an upper-bound because the 2 coefficients are zero for even continuity orders (see Fig. 22) and has been neglected from the equation. The result of Fig. 17 can also be deduced by comparing twice Eq. 27 to Eq. 28. Additionally, the memory overhead of 1D orthogonality is abated by realising that the higher-order continuity requires that the lower-order B-coefficients already be read from main memory and are therefore already stored in a register. The DAB algorithm converges in $\hat{r} \cdot n$ iterations with the algorithm's parallel arithmetic complexity given by Eq. 29. Both the original and the orthogonal formulation only need to project until convergence in each direction once (i.e. applying continuity in one direction does not impact continuity in the previous directions). It is presumed to be due to the layers being orthogonal to each other; however, further analysis is necessary.




Fig. 19 The original continuity equations as stated by De Visser [35] and Theorem 3. The B-coefficients that are affected by the continuity equations are filled.




Fig. 20 The original continuity equations orthogonalised to all preceding continuity orders. The B-coefficients that are affected by the continuity equations are filled.




Fig. 21 The original continuity equations as stated in [35] and Theorem 3, for seventh order continuity. The B-coefficients that are affected by the continuity equations are white.




Fig. 22 The original continuity equations orthogonalised to all preceding continuity orders, for seventh order continuity. The B-coefficients that are affected by the continuity equations are white.

$$O(\text{Orig/iteration}) = \sum_{m=0}^{\hat{r}} (m+2) = \frac{\hat{r}^2 + \hat{r}}{2} + 2(\hat{r}+1) \quad (27)$$

$$O(\text{Orth}) < \sum_{m=0}^{\hat{r}} 2(m+1) = \hat{r}^2 + \hat{r} + 2(\hat{r}+1) \quad (28)$$

$$O(\text{DAB}) \approx O(\hat{r}^2 n) \quad (29)$$

7 Memory Layout

This section will describe the two modifications required for better use of the memory of a system to ensure parallel execution of the algorithm.

7.1 Morton Z-ordering

Morton Z-ordering is a space-filling curve that is often used for organising data in GPUs to optimise accesses for spatial locality. The main aim of this subsection is to mitigate cache issues such as false sharing and cache pollution, as mentioned in Section 3.3. Eliminating these issues is critical for using the cache hierarchy effectively and is beneficial on both CPUs and GPUs.

With orthotope B-splines, there are 3 main z-orderings: orthotope-centred degree-based (OD) ordering, orthotope-centred continuity-based (OC) ordering, and vertex-centred continuity-based (VC) ordering. Orthotope-centred means all coefficients belonging to a orthotope are stored in a contiguous section of memory. In contrast,




Fig. 23 The indexing starts at the blue dot and follows the arrows. This figure illustrates orthotope-centred degree-based (OD) ordering in 2D.




Fig. 24 The indexing starts at the blue dot and follows the arrows. This figure illustrates orthotope-centred continuity-based (OC) ordering in 2D.

vertex-centred stores all coefficients belonging to a vertex in contiguous memory. Vertex-centred ordering has the potential to save up to $2^n - 1$ global memory transactions when applying the constraints, where n is the number of dimensions. However, this benefit is mitigated by the decrease in performance when determining the unconstrained optimal solution and when evaluating the spline. Due to the resorting, extracting the data for a given spline becomes more complicated as the data is no longer contiguous. The benefit of vertex-centred ordering is further diminished when using shared memory on the GPU, as the read from global memory only needs to be performed once per kernel call. OD ordering is ordering the coefficients based on the multinomial coefficients of the spline and is identical to the ordering presented in [35] (see Fig. 23). OC ordering begins at each vertex and then moves toward the centre of the orthotope (see Fig. 24). All coefficients belonging to a vertex of an orthotope are stored contiguously improving cache performance (less cache pollution) on both the CPU and GPU when applying the constraints, enabling better memory access coalescing, and simplifying the indexing function. OC is expected to outperform OD because false sharing is avoided - as the cache lines are no longer shared between two vertices in an orthotope. False sharing is particularly detrimental to performance as the data needs to be read and written to after each projection. No quantitative analysis of the impact of false sharing on performance is available, as a full implementation of the algorithm has yet to be made. It should be noted that an additional padding may be necessary for the OC ordering to ensure two neighbouring vertices within an orthotope do not share cache lines. A non-performance benefit is that the indexing function is symmetric about its edges and means offsets are always positive away from the edge. The offset for each direction is always a power of the number of continuity equations \hat{r} .

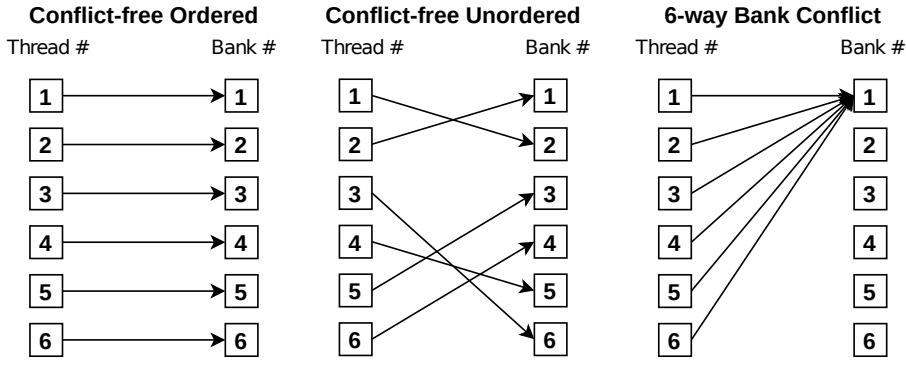


Fig. 25 Each thread accesses a sequential memory bank.

Fig. 26 Each thread accesses a different memory bank. There is no performance benefit compared to Fig. 25.

Fig. 27 All threads access the same memory bank. If the accesses are to different addresses, the accesses are performed sequentially.

7.2 Bank Conflict Avoidance

An additional memory consideration for GPUs, Field-programmable Gate Arrays (FPGAs), Application Specific Integrated Circuits (ASICs), and parallel hardware in general is that the memory can be split into multiple separate chips and is often referred to as memory banks. Each memory bank can be accessed in parallel enabling a higher data throughput. Additionally, it ensures that parallel threads do not need to stall for sequential accesses to memory. Typically, the data in the memory banks is striped, similar to RAID 0, meaning sequential memory addresses are mapped to different memory banks. Modern Nvidia GPUs, work with 32 threads in parallel and have 32 memory banks and, ideally, that means each thread has its own memory bank; however, as the data is striped, it is possible that two or more threads need to access the same memory bank and this is referred to as a bank conflict (see Fig. 27). Fine-grained control on avoiding bank conflicts is only provided by using shared memory on Nvidia GPUs and has the additional benefit of reducing memory bandwidth contention⁵ when performing writes after each projection step of the algorithm. The total amount of shared memory required to store all the data for a vertex for different numbers of dimensions and continuities is given in Table 1. A first-order model in 7 dimensions is the limit for shared memory in modern hardware. B-coefficient deduplication, as in Appendix B, has the potential to increase this limit, if a conflict-free memory ordering is possible.

Bank conflict avoidance is a graph colouring problem, where the number of permitted colours is determined by the number of banks. For multivariate orthotope B-splines, the B-coefficients are the vertices of the graph and all coplanar coefficients on axis-aligned hyperplanes are connected by edges. Mathematically, an axis-aligned hyperplane is any hyperplane that changes in at most $n - 1$ directions. The problem is similar in some respects to the sudoku graph colouring problem with an additional

⁵ This is due to the write-through cache policy.

Table 1 The number of coefficients that are contained within a vertex for different continuities and numbers of dimensions. The amount of shared memory required is the entry multiplied by the number of bytes for the data type. A 7th dimensional orthotope with first-order continuity requires 64 KiB of shared memory when using single-precision floating points and would just fit in shared memory.

Dimensions Continuity	1	2	3	4	5	6	7	8
0	2	4	8	16	32	64	128	256
1	4	16	64	256	1024	4096	16384	65536
2	6	36	216	1296	7776	46656	279936	1679616
3	8	64	512	4096	32768	262144	2097152	16777216




Fig. 28 A 4D hypercube can be interpreted as two 3D hypercubes (cubes) connected at the same vertices. The 3D hypercubes can then be decomposed into a pair of 2D hypercubes (squares).

relaxation that colours may reoccur within a hyperplane provided no bank conflict occurs (proper colouring). This also relaxes the requirement of connecting all vertices on the same hyperplane to a subset of the vertices, where the number of connected vertices is determined by the number of banks available. Alternatively, when this is not possible, the number of bank conflicts should be minimised (improper colouring), which is an optimisation problem in itself. Moreover, the colours (or indices) should be computationally-efficient to compute to mitigate the cost of bank conflicts. First-order continuity for a trivariate regression problem requires $3^3 = 27$ threads and $3^3 + 3^2 = 36$ banks for a simple algorithm for conflict-free access to a block of data (see Fig. 29). This leaves 5 threads idle but results in more constraints being computed per iteration. Higher than first-order continuity with B-coefficient deduplication cannot be efficiently solved with this approach as it does not generalise to higher continuities without requiring a significant increase in the number of memory banks. Instead, without B-coefficient deduplication, it is possible to create a (2,2,2)-cube of B-coefficients that only requires 8 banks. This conflict-free block can then be tiled in higher dimensions guaranteeing conflict-free access for all out-of-edge dimensions. Additionally, since 8 banks are required for the cube, all 32 banks and threads can be used when 8 cubes are available - which is the case for first-order continuity in 4 dimensions.




Fig. 29 One plausible conflict-free tiling that can be used for (3,3,3)-cube; however, it requires 36 memory banks and Nvidia GPUs only have 32. Additionally, only 27 threads will typically be active instead of the maximum capacity of 32, although this fact is outweighed by reducing redundant computation by using deduplication (see Appendix B). The amount of memory padding required can be mitigated using modulo addressing, as each address in a (D1,D2,D3)-cube is only used once.




Fig. 30 This figure illustrates which B-coefficients needs to be stored in each memory bank and the order in which they should be read or written to ensure that accesses are conflict-free when performing Kaczmarz's algorithm. The algorithm distinguishes access patterns across constraints in the directions D1, D2, and D3 to those in D4 and higher. For D4 and higher, the B-coefficients enclosed in a square (e.g. blue square) can be assigned to a single warp. For D1, the B-coefficients read and written to by a single warp is coloured in red. The approaches for D2 and D3 follow analogously from D1. Notice that the vertex number still indicate V1-V8 when moving in the D4 direction, this is because D4 is four copies of the (D1,D2,D3)-cube wide and only two copies have been shown for readability.

8 Conclusions & Recommendations

Orthotope B-splines are better for high-dimensional regression as the algorithms required to implement them are simpler and generalise better than those for high-dimensional simplex splines. In this paper, oblique projections and first-order optimisation were demonstrated to improve results by avoiding overfitting - which is often the case in high-dimensional space - and avoiding constraints to be applied between splines when one spline is missing data, as it reduced the training and validation accuracy. Moreover, orthotope B-splines with a different memory organisation has

the potential to significantly improve performance, on both a CPU and a GPU as the cache hierarchy and memory bandwidth can be better utilised. Additionally, the (2,2,2)-cube tiling is compact and results in no padding of shared memory. Finally, a conflict-free use of shared memory on the GPU is critical to enable the scalability and parallelisability of the algorithm.

Further research should investigate using directional derivatives to improve the model in regions where insufficient data is available - while avoiding overfitting. Additionally, investigating how this influences model accuracy in the online setting. It was noted the first projection direction biases the constrained solution to favour the initial projection directions over the latter directions. Research into a weighted solution that minimises this bias should be evaluated to see if it influences model accuracy. A CUDA implementation can be useful to run larger models and enable online experiments to be performed to evaluate the algorithm in situ. The current algorithm is primarily suited for first-order continuity. Further research could look at memory tilings for higher-order continuity, which may be beneficial in other fields like structural analysis. Finally, hybridised optimisation approaches, where the Kaczmarz method is used as a smoothing filter to ensure constraints are satisfied if, for example, dual ascent is terminated before convergence of the dual function.

References

1. van den Aarsen, M.: Distributed approach for aerodynamic model identification of the ice aircraft. Master's thesis, Delft University of Technology (2018). URL <http://resolver.tudelft.nl/uuid:df32613c-5f3c-484e-84ea-99672939d6a5>
2. Awanou, G., Lai, M.J., Wenston, P.: The multivariate spline method for scattered data fitting and numerical solutions of partial differential equations. In: G. Chen, M.J. Lai (eds.) *Wavelets and Splines*: Athens 2005, pp. 24–75. Nashboro Press (2005)
3. Bertsekas, D.P.: *Convex Optimization Algorithms*, first edn. Optimization and Computation Series. Athena Scientific (2015)
4. Bin Zubair, H.: Efficient multigrid methods based on improved coarse grid correction techniques. Ph.D. thesis, Delft University of Technology (2009). URL <http://resolver.tudelft.nl/uuid:e0b43b38-ad07-4076-baf1-aa02579c397f>
5. Boeing: Statistical summary of commercial jet airplane accidents: Worldwide operations 1959-2018. Online (2019). 50th Edition
6. Boyd, S., Parikh, N., Chu, E., Peleato, B., Eckstein, J.: Distributed optimization and statistical learning via the alternating direction method of multipliers. *Foundations and Trends in Machine Learning* **3**(1), 1–122 (2010)
7. Chen, C., He, B., Ye, Y., Yuan, X.: The direct extension of admm for multi-block convex minimization problems is not necessarily convergent. *Mathematical Programming* **155**(1-2), 57–79 (2014). DOI 10.1007/s10107-014-0826-5
8. de Visser, C.C., Brunner, E., Verhaegen, M.: On distributed wavefront reconstruction for large-scale adaptive optics systems. *Journal of the Optical Society of America A* **33**(5), 817–831 (2016). DOI 10.1364/JOSAA.33.000817
9. de Visser, C.C., Verhaegen, M.: Wavefront reconstruction in adaptive optics systems using nonlinear multivariate splines. *Journal of the Optical Society of America A* **30**(1), 82–95 (2013). DOI 10.1364/josaa.30.000082
10. Duchi, J., Hazan, E., Singer, Y.: Adaptive subgradient methods for online learning and stochastic optimization. *Journal of Machine Learning Research* **12**, 2121–2159 (2011)
11. Dykstra, R.L.: An algorithm for restricted least squares regression. *Journal of the American Statistical Association* **78**(384), 837–842 (1983). DOI 10.2307/2288193
12. Dykstra, R.L., Boyle, J.P.: An algorithm for least squares projections onto the intersection of translated, convex cones. *Journal of Statistical Planning and Inference* **15**, 391–399 (1986-1987). DOI 10.1016/0378-3758(86)90111-4

13. European Aviation Safety Agency: Executive director decision 2015/012/r. Online (2015)
14. European Aviation Safety Agency: Explanatory note to decision 2015/012/r. Online (2015)
15. Federal Aviation Administration: Qualification, service, and use of crewmembers and aircraft dispatchers. Online (2013). URL <https://www.govinfo.gov/content/pkg/FR-2013-11-12/pdf/2013-26845.pdf>
16. Gaffke, N., Mathar, R.: A cyclic projection algorithm via duality. *Metrika* **36**, 29–54 (1989). DOI 10.1007/BF02614077
17. Goodfellow, I., Bengio, Y., Courville, A.: *Deep Learning*. MIT Press (2016). <http://www.deeplearningbook.org>
18. Hooij, B.: Online distributed approach for aerodynamic model identification of the ice aircraft. Master's thesis, Delft University of Technology (2020). URL <http://resolver.tudelft.nl/uuid:a92817a5-b7f9-4d5f-84de-30f6bd78f668>
19. Huisman, F.J.: Full flight envelope aerodynamic modelling of the cessna citation ii using physical splines. Master's thesis, Delft University of Technology (2017). URL <http://resolver.tudelft.nl/uuid:f9926c24-e86b-4eee-8986-d427fb4bc667>
20. International Air Transport Association: Guidance Material and Best Practices for the Implementation of Upset Prevention and Recovery Training, second edn. International Air Transport Association (2018)
21. International Air Transport Association: Loss of Control In-Flight Accident Analysis Report. International Air Transport Association (2019)
22. Jacklin, S.A.: Closing the certification gaps in adaptive flight control software. In: AIAA Guidance, Navigation and Control Conference and Exhibit. American Institute of Aeronautics and Astronautics (2008). URL <https://ntrs.nasa.gov/archive/nasa/casi.ntrs.nasa.gov/20090026333.pdf>
23. Kaczmarz, S.: Angenäherte auflösung von systemen linearer gleichungen. *Bull. Int. Acad. Polon. Sci. A*. pp. 355–357 (1937)
24. Karagöz, R.: Tensor network b-splines for high-dimensional function approximation. Master's thesis, Delft University of Technology (2020). URL <http://resolver.tudelft.nl/uuid:d637a2f4-512d-4330-83ae-29f7ef2958ed>
25. Keskar, N.S., Socher, R.: Improving generalization performance by switching from adam to sgd. Tech. rep., Salesforce Research (2017). URL <https://arxiv.org/abs/1712.07628.pdf>
26. Kingma, D.P., Ba, J.: Adam: A method for stochastic optimization. In: 3rd International Conference for Learning Representations (2015). URL <https://arxiv.org/abs/1412.6980>
27. Krizhevsky, A., Sutskever, I., Hinton, G.E.: Imagenet classification with deep convolutional neural networks. In: *Advances in neural information processing systems*, pp. 1097–1105 (2012)
28. Lévy, J., Ren, S., Mushtaq, H., Bertels, K., Al-Ars, Z.: GASAL2: a GPU accelerated sequence alignment library for high-throughput NGS data. *BMC Bioinformatics* (2019). DOI <https://doi.org/10.1186/s12859-019-3086-9>
29. McCloskey, M., Cohen, N.J.: Catastrophic interference in connectionist networks: The sequential learning problem. *Psychology of Learning and Motivation* **24**, 109–165 (1989). DOI 10.1016/S0079-7421(08)60536-8
30. Mnih, V., Kavukcuoglu, K., Silver, D., Rusu, A.A., Veness, J., Bellemare, M.G., Graves, A., Riedmiller, M., Fidjeland, A.K., Ostrovski, G., Petersen, S., Beattie, C., Sadik, A., Antonoglou, I., King, H., Kumaran, D., Wierstra, D., Legg, S., Hassabis, D.: Human-level control through deep reinforcement learning. *Nature* **518**, 529–533 (2015). DOI 10.1038/nature14236
31. Nesterov, Y.: A method for unconstrained convex minimization problem with the rate of convergence $O(1/k^2)$. *Doklady AN USSR* **269**, 543–547 (1983)
32. Shewchuk, J.R.: An introduction to the conjugate gradient method without the agonizing pain. Tech. Rep. 1.25, Carnegie Mellon University (1994)
33. Smaragdous, G., Chatzikonstantis, G., Kukreja, R., Sidiropoulos, H., Rodopoulos, D., Sourdis, I., Al-Ars, Z., Kachris, C., Soudris, D., Zeeuw, C.I.D.: Brainframe: a node-level heterogeneous accelerator platform for neuron simulations. *Journal of Neural Engineering* (2017). DOI <https://doi.org/10.1088/1741-2552/aa7fc5>
34. van der Peijl, I.: Physical splines for aerodynamic modelling of innovative control effectors. Master's thesis, Delft University of Technology (2017). URL <http://resolver.tudelft.nl/uuid:7730dd69-2bb8-4047-8f6d-5681cfd04cee>
35. de Visser, C.C.: Global nonlinear model identification with multivariate splines. Ph.D. thesis, Delft University of Technology (2011). URL <http://resolver.tudelft.nl/uuid:6bc0134a-0715-4829-903d-6479c5735913>

36. Visser, T., de Visser, C., van Kampen, E.: Quadrotor system identification using the multivariate multiplex b-spline. In: AIAA Atmospheric Flight Mechanics Conference. AIAA (2015). DOI doi.org/10.2514/6.2015-0747. URL <http://resolver.tudelft.nl/uuid:6e50cda5-3c14-415e-b800-3e1dd42102b5>
37. Wei, E., Ozdaglar, A.: Distributed alternating direction method of multipliers. In: 2012 IEEE 51st IEEE Conference on Decision and Control, pp. 5445–5450. IEEE (2012). DOI [10.1109/CDC.2012.6425904](https://doi.org/10.1109/CDC.2012.6425904)
38. Wei, E., Ozdaglar, A.: On the $O(1/k)$ convergence of asynchronous distributed alternating direction method of multipliers. In: 2013 IEEE Global Conference on Signal and Information Processing, pp. 551–554. IEEE (2013). DOI [10.1109/GlobalSIP.2013.6736937](https://doi.org/10.1109/GlobalSIP.2013.6736937)
39. Zhu, Y., Rong Li, X.: Recursive least squares with linear constraints. Communications in Information and Systems **7**(3), 287–312 (2007). URL <https://projecteuclid.org/euclid.cis/1201531976>

A B-spline Directional Derivatives

Directional derivatives are useful for analysis, especially in the context of aerodynamic system identification and control. The directional derivatives in physical space can, for example, be used to assess the static stability of the aircraft. They can also be used to design a controller to meet certain requirements or quantify the inability to meet them.

The De Casteljau Matrix from [35] - given in Eq. 31, where each entry is formed using Eq. 30 - is a reformulation of the recursive De Casteljau algorithm into a non-recursive matrix form. \hat{d}_1 and \hat{d}_2 are given by filling in the first and second superscripts into Eq. 4, respectively.

$$P_{\gamma}^m(b) = \begin{cases} \frac{m!}{\gamma!} b^m & \text{if } \gamma \geq 0 \\ 0 & \text{otherwise} \end{cases} \quad (30)$$

$$\mathbf{P}^{d,d-m}(b) = \left[P_{\theta-\kappa}^m(b) \right]_{\substack{\theta \in \mathbb{M}_d^{n+1} \\ \kappa \in \mathbb{M}_{d-m}^{n+1}}} \in \mathbb{R}^{\hat{d}_2 \times \hat{d}_1} \quad (31)$$

The directional derivative is taken with respect to a directional vector u . Unfortunately, u is a vector and cannot be converted to barycentric coordinates; however, to overcome this problem, the vector u can be described as the difference between two points in Cartesian space (see Eq. 34) and these points can then be transformed into barycentric space and then subtracted from each other and denoted as a (see Eq. 35) [35]. The m^{th} directional derivative of the polynomial p with respect to the direction u is given by Eq. 36.

$$B_{\kappa}^d(b) = \sum_{|\gamma|=m} P_{\gamma}^m(b) B_{\kappa-\gamma}^{d-m}(b) \quad (32)$$

$$p(b) = \mathbf{B}^{d-m}(b) \mathbf{P}^{d,d-m}(b) \cdot c^t j \quad (33)$$

$$u = v - w \in \mathbb{R}^n \quad (34)$$

$$b(u) := a = b(v) - b(w) \in \mathbb{R}^{n+1} \quad (35)$$

$$D_u^m p(b) = \frac{d!}{(d-m)!} B^{d-m}(b) \mathbf{P}^{d,d-m}(a) \cdot c^t j \quad (36)$$

B B-coefficient Deduplication

This section has been placed in the appendix as the B-coefficient deduplication does not work well in combination with Section 7.2. This is because the deduplication leads to odd widths for the vertex blocks and the bank conflict avoidance scheme requires tilings of even width blocks due to the limited number of memory banks.

At the edges connecting two splines, the unconstrained B-coefficients represent the same physical-space location but may have different values. After applying zeroth-order continuity across the edge, they are also equal in value. If these duplicated coefficients are not removed, each iteration will require that

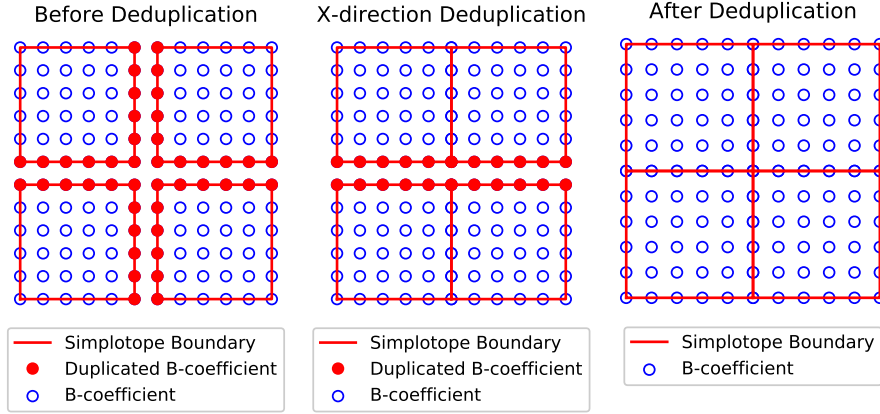


Fig. 31 Splines have been artificially separated to illustrate the duplicated coefficients.

Fig. 32 B-coefficients in the x-direction have been deduplicated by applying zeroth-order continuity in the x-direction. Splines have been artificially separated to illustrate the duplicated coefficients.

Fig. 33 B-coefficients in the y-direction have been deduplicated by applying zeroth-order continuity in the y-direction. Splines have been artificially separated to illustrate the duplicated coefficients.

zeroth-order continuity be reapplied and lead to slower overall convergence. This is obviated by using orthogonal constraints, as mentioned in Section 6; however, this still requires additional memory store operations and multiplications - if not deduplicated. Deduplication improves scalability as base of the exponential growth in higher dimensions is reduced by one, as given by Eq. 37. For first-order continuity in the 3D case, 2.37 times fewer memory is required and the benefit grows with more dimensions. The amount of redundant computation is significantly reduced as one fewer multiplication is required per projection and the number of required projections decreases from $(2 \cdot \hat{r})^{n-1}$ to $(2 \cdot \hat{r} - 1)^{n-1}$. A lower memory load means more data can be fit into the cache or, in the case of CUDA, shared memory. The deduplication is performed using twice the duplicated coefficient's value in Eq. 25 and updating the coefficient using Eq. 26 with half its value.

$$\mathbf{M} = \frac{(2 \cdot \hat{r})^n}{(2 \cdot \hat{r} - 1)^n} \quad (37)$$

When utilising oblique projections, the B-coefficient deduplication process needs to be weighted based on the number of splines with useful data. The weighting is changed each iteration to be the sum of the weights of the two splines involved for that direction. Each orthotope needs a weight for each direction. The weight for a B-coefficient is determined as the sum of the weights of the directions it is involved with. For example in the 2D case, a vertex of an orthotope is involved in 2 directions and an edge is only involved in one direction.

C Performance Analysis

This section will analyse the speedup of the algorithm for modifications to the storage of data. The algorithms have been implemented in *MATLAB* and the results may differ when optimising the algorithms for other platforms or in other programming languages. The DAB algorithm can operate in two main scenarios: online and batch. When using the algorithm in the online regression setting, typically the computational capabilities is significantly decreased as in the case of micro unmanned aerial vehicles, where the L1 cache size is typically around 8 KB.

The experiments were conducted on a hexacore Ryzen 3600 processor with a clock frequency of 4.2 GHz and 32 GB of RAM at 2933 MHz. The processor has 32 KB of L1 data cache per core and 32 MB




Fig. 34 The speedup over using the *MATLAB* sparse matrix format for the constraint matrix H in the 3D case with different resolution grids. The benefit of on-the-fly and the proposed sparse format improves as the grid becomes finer.

of shared L3 cache. The data for approximately 500 single-precision floating-point 7D orthotopes could in theory fit in the L3 cache. Of course, in practice the number of orthotopes will be slightly less as no fine-grained control of the cache is available. Additionally, if a sparse matrix format is used - in contrast to an on-the-fly approach - this data will also consume cache space. L3 cache is; however, 10 times slower than L1 cache with respect to latency on the selected system. Execution times are averaged over 10 runs and the code tested is, specifically, the orthogonal Kaczmarz projection algorithm with orthogonal constraints.

One optimisation applied to the DAB algorithm was to reduce the total amount of memory required. The constraint matrix is very sparse. Using a sparse matrix format requires storing the indices and the values of each non-zero entry. The values are constant for a given continuity and can be stored separately from the indexing. The indexing data can be stored in an array with a column for each non-zero entry in a row and a row for each constraint. The columns of the Index array maps one-to-one with the columns of the Value array. There is one Value and one Index array per continuity order and per continuity direction as the dimensions of the arrays are dependent on the continuity order and each direction is handled separately. The proposed approach is similar to the Compressed Sparse Row (CSR) format except with a set of values per continuity order instead of per non-zero entry and the row pointer is implicit as all rows have the same width. Finally, a comparison will be performed with on-the-fly computation of the indices, as computing the indices and the neighbouring orthotopes is trivial with the regularity of the tessellation. It should be noted that *MATLAB* uses the Compressed Sparse Column (CSC) format for storing sparse matrices. In Fig. 34, the *MATLAB* sparse storage format outperforms the on-the-fly approach for small tessellations; however, as the size of the tessellation increases, so does the speedup of the on-the-fly approach. From Fig. 34, Fig. 35, and Fig. 36, it can be seen that the proposed sparse format always outperforms the sparse method implemented by *MATLAB*. With increased dimensions and tessellation sizes, the proposed sparse format is the best up to at least the 5D case, as tested; however, a slight decrease in performance is observed in Fig. 35 and Fig. 36, as the L3 cache is filled beyond its capacity. The benefit of the proposed sparse method is also more aimed at CPU computation, as GPUs are often memory bottlenecked and the larger number of cores are better suited to on-the-fly computation. The execution times are provided in Table 2. It should be noted that all implementations utilised the optimisation that the denominator of the projection is the same for all projections for a given continuity order and does not require an expensive dot product.




Fig. 35 The speedup over using the *MATLAB* sparse matrix format for the constraint matrix H in the 4D case with different resolution grids. The proposed sparse method has a reduced speedup for the case of 8 orthotopes per dimension because the limits of the L3 cache have been reached. The on-the-fly approach still improves as computing the values is faster than loading them from main memory.




Fig. 36 The speedup over using the *MATLAB* sparse matrix format for the constraint matrix H in the 5D case with different resolution grids. The on-the-fly approach also begins to lose performance as the computed indices are being temporarily sent to main memory. Decomposing the on-the-fly computation into batches may further improve the performance.

Table 2 This table shows the average execution times in seconds for different tessellation sizes using the different approaches. In the 5D case with 8 orthotopes per dimension, the execution time takes around 5.2 seconds to complete with the fastest method.

Dimensions	Sparse (MATLAB)			Sparse (Proposed)			On-the-fly		
	2/dim	4/dim	8/dim	2/dim	4/dim	8/dim	2/dim	4/dim	8/dim
3D	3.05E-04	5.74E-04	0.0031	2.81E-04	4.05E-04	0.0017	6.65E-04	8.11E-04	0.0027
4D	6.68E-04	0.0076	0.1787	4.04E-04	0.0028	0.078	9.39E-04	0.004	0.0858
5D	0.0047	0.2492	14.0383	0.0015	0.0769	5.2348	0.0026	0.0877	5.5535